

1 Teaching

I enjoy teaching as a way of having a positive effect on how people think and as a way of refreshing and sharpening my own understanding of concepts and tools. I have taught Econometrics I three times in Chile and two semester-long courses at UC Berkeley (Introduction to Econometrics and Introduction to Environmental Economics and Policy), alongside multiple teaching assistantships and tutoring roles. I have also led dozens of research transparency workshops in several countries.

Teaching Philosophy: Clarity and Rigor

In both my research and my teaching, I aim to raise the floor and not only focus on pushing the frontier. By raising the floor I mean helping every student along, not just the brightest ones, and making the average policy analysis more credible, not just the best one. Transparency in policy analysis and in reporting raises the floor in research. Clarity and rigor raise it in teaching. As an example of clarity, in previous courses I have spent significant time discussing multiple ways to grasp the Central Limit Theorem, rather than focusing only on its formal proof. Rigor starts with teaching materials that are computationally reproducible. Students can see how every figure and result is derived, and can change any input to see how the result moves. My materials for [Introduction to Econometrics](#), which I taught at UC Berkeley in 2022, focused on both clarity and rigor to increase understanding and raise the floor.

My approach to teaching is also designed to help make academic institutions of excellence more accessible to students from disadvantaged socioeconomic backgrounds. Here raising the floor means that I focus on students in the left tail of the distribution and encourage them to set realistic goals, such as substantial improvements relative to their starting point. This pushes against the “meritocratic consensus” documented by Bleemer and Rothstein (2026). In my last course, Introduction to Econometrics, I encouraged students throughout the term to use office hours, and I approached individual students to urge them to come. In that course I found that disadvantaged students did not know about additional campus resources such as tutoring and advising. Going forward I would help ensure students know about and can access those resources the same way: list them in the syllabus, then keep encouraging students, throughout the semester, to use them with confidence and ownership.

Courses

I am well-prepared to teach methods for policy analysis, program evaluation, and applied econometrics, as well as most undergraduate courses in economics and statistics. Lee Friedman’s work has influenced my thinking on policy analysis (Friedman, 2017). A course on Open Policy Analysis (see my research statement) is one I would especially like to teach, and I have written a [course proposal](#) for it. I would build assignments around real policy questions, because students engage more with a method once they see it solve a problem they care about. Students would use artificial intelligence (AI) coding agents to build a transparent and reproducible analysis of a current policy question, such as a proposed wealth tax, and would verify every step. They would leave with a portfolio piece that any reader can inspect and rerun.

Teaching and Learning with AI

In mathematics, agentic AI now generates proofs faster than mathematicians can “digest” them (Tao, 2026). Digesting a proof means clarifying and explaining it to other mathematicians and to other disciplines, so that they can build applications and new knowledge on top of it. The lesson I draw for teaching is that producing gradable output is no longer the challenge. Understanding it is. I plan to teach students to use AI to produce high-quality output, then to understand, explain, and modify that output, and to feel confident disclosing their use of AI. One sentence I keep going back to, popularized by Andrej Karpathy (Karpathy, 2026), is “You can outsource your thinking, but you cannot outsource your understanding.”

I now use agentic harnesses, such as Claude Code, heavily in my day-to-day work. A harness turns an AI model into a coding agent that runs code and carries out multi-step tasks on a user's computer. For example, over the past year I have used them to build short course materials for my own learning, such as a [refresher on machine learning](#) framed around algorithmic discrimination in ride-sharing platforms. Each of its five modules has slides, written notes, and an R exercise, and the current version took me less than a week to develop.

My courses would steer students toward using these tools to learn, test, and expand their understanding, and away from output they neither own nor understand. Students need a basic understanding of agentic harnesses. Everyone who uses them should learn concepts once left to programmers: version control, Markdown documentation, verification plans, token management, and global and local memory. In each course that uses these tools, I plan to spend the first couple of classes on these concepts.

I plan to assess what a student understands, not how finished the work looks. Verification is the first skill I will teach. I will give students AI-generated output, tell them that one section contains hallucinations or errors, and ask them to find the errors. Explanation is the test of understanding. To improve and measure it, I will combine traditional tools, such as in-class tests, with AI tools, such as a routine that questions a student one step at a time until every part of the reasoning is explicit. Disclosure of AI use is harder to secure. I will lead by example, as I do in my own [open policy analysis of the economic effects of AI](#), and explore incentive-compatible mechanisms that help students disclose their use of AI.

2 Mentoring

Between 2018 and 2022 I supervised 27 Berkeley undergraduates through the Undergraduate Research Apprentice Program (URAP). Since 2019 I have supervised eight research assistants (RAs). My former apprentices and assistants have gone on to academic institutions and to public and private organizations. My CV lists each placement.

URAP apprentices need a bounded task that they can finish in at most 10 weeks at 10 hours of work per week. RAs need clear but broad instructions, with an agreement at the start on the scope of the work and on how it will be credited. I plan to train both groups to verify AI output and to disclose their use of it. As a faculty member I would advise PhD students, serve on dissertation committees, and supervise capstone projects, helping students scope an analysis that a client can act on and any reader can inspect.

3 Service: Reviewing, Public Goods, and Networks

I have refereed for the *Harvard Data Science Review* and *Contemporary Economic Policy*. As a faculty member I would take on the committee work expected of junior faculty, and I can contribute most where reproducible analysis or the use of AI tools in policy work is at issue.

An important part of my service would be producing public goods. I co-designed and oversaw the [Social Science Reproduction Platform](#), a tool for crowdsourcing computational reproducibility, which hosted over 300 reproductions of more than 200 papers.¹ I have also produced several open policy analyses, and I plan to keep expanding their number. Those I build in my research, and those students build in class, can serve policy analysis organizations at the governmental level (e.g., the Congressional Budget Office and California's Legislative Analyst's Office), the non-governmental level (e.g., GiveWell), and the multilateral level (e.g., the World Bank).

Finally, my networks in Latin America, Norway, and across UC Berkeley, including the Center for Effective Global Action (CEGA), can support student and scholar exchanges and connect my future colleagues and students with other departments and centers.

¹Archived version. Its companion [guide to computational reproducibility](#) remains online.

References

Zachary Bleemer and Jesse Rothstein. The meritocratic consensus and stratification in higher education. *Journal of Economic Perspectives*, 2026. Forthcoming.

Lee S. Friedman, editor. *Does Policy Analysis Matter? Exploring Its Effectiveness in Theory and Practice*. Wildavsky Forum Series. University of California Press, 2017.

Andrej Karpathy. Sequoia Ascent 2026 summary. Blog post, April 30, 2026.

Terence Tao. Mathematics in the age of AI. arXiv:2608.16753, 2026.